

# Project Portfolio Ideas for FAANG Aspirants: A Guide for Mid-Senior Java/Docker Developers

## 1. Introduction

Securing a mid-senior level Java/Docker developer role at a FAANG (Facebook, Amazon, Apple, Netflix, Google) company represents a significant career milestone. These organizations operate at the forefront of technological innovation, tackling complex challenges that demand a high degree of expertise and practical experience. In such a competitive landscape, a well-crafted GitHub portfolio serves as a powerful tool for developers to distinguish themselves. It provides tangible evidence of their capabilities, offering recruiters and hiring managers a window into their practical skills and project experience that goes beyond a traditional resume. This report aims to provide concrete and compelling project ideas, accompanied by detailed implementation outlines, specifically tailored to showcase the skills and knowledge highly valued by FAANG companies. The focus will be on projects that demonstrate proficiency in areas such as microservices architecture, distributed systems, and cloud-native technologies, all critical components of modern enterprise-level software development.

## 2. What FAANG Looks For in Mid-Senior Java/Docker Developers

FAANG companies prioritize candidates who possess a strong theoretical foundation coupled with demonstrable practical experience in building and scaling complex software systems. They seek developers who can not only write efficient and robust code but also understand the architectural considerations involved in designing and deploying large-scale applications. Several key technical skills and project characteristics consistently capture the attention of FAANG recruiters and hiring managers.

### Microservices Architecture

The microservices architectural style has become a dominant pattern in modern enterprise application development. This approach, where an application is structured as a collection of small, independent services communicating over a network, offers numerous benefits in terms of scalability, independent deployments, and fault isolation. The prevalence of microservices is evident in the abundance of open-source projects and extensive discussions surrounding their implementation using Java and Docker.<sup>1</sup> FAANG companies frequently adopt microservices to manage the complexity and scale of their offerings, allowing different teams to work independently on specific functionalities. Therefore, demonstrating the ability to design, build, and deploy microservices using Java and Docker is a significant advantage for candidates. The extensive material available on this topic suggests that it is a core area of interest and a likely requirement for engineers working on large-scale, distributed systems within these organizations.

### Distributed Systems

Operating at a massive scale necessitates a deep understanding of distributed systems concepts. This includes familiarity with distributed databases, such as Apache HBase and Apache Cassandra <sup>22</sup>, distributed task queues, and distributed coordination services like Apache Zookeeper and etcd.<sup>22</sup> Projects like Apache Airavata, a software framework for

executing and managing computational jobs on distributed computing resources <sup>24</sup>, further highlight the relevance of distributed computing frameworks in the Java ecosystem. FAANG companies' systems are inherently distributed to handle the immense volume of data and user traffic they experience. Consequently, developers with a strong grasp of distributed systems principles, including concepts like consistency, availability, and fault tolerance, are highly sought after. The ability to design and implement systems that can maintain data integrity across multiple nodes, handle high loads efficiently, and gracefully recover from failures is crucial for engineers in these environments. Projects demonstrating these skills provide compelling evidence of a candidate's readiness for such challenges.

### **Containerization with Docker**

Docker has become a fundamental technology for packaging and deploying applications, especially in the context of microservices and cloud-native environments. Its importance is underscored by its frequent mention across various resources related to Java, microservices, and cloud computing.<sup>1</sup> Proficiency in Docker is often considered a baseline expectation for mid-senior developers. Portfolio projects should not only demonstrate basic Docker usage but also showcase best practices in Dockerfile creation, such as optimizing image size through multi-stage builds <sup>35</sup>, and efficient container management. Docker provides isolation, portability, and scalability, making it an indispensable tool for FAANG companies that heavily rely on containerized deployments to ensure consistency and efficiency across their infrastructure.

### **Container Orchestration (Kubernetes/Docker Swarm)**

Managing and scaling containerized applications effectively requires the use of container orchestration tools. Kubernetes and Docker Swarm are the leading technologies in this domain.<sup>4</sup> Experience with these tools demonstrates a developer's ability to manage large-scale deployments, handle scaling demands, and ensure high availability of applications. FAANG companies manage a multitude of services deployed as containers, and familiarity with orchestration platforms like Kubernetes or Docker Swarm is a significant asset, showcasing the capacity to handle production-level complexity and maintain system reliability.

### **Cloud-Native Development (AWS, GCP, Azure)**

The increasing reliance on cloud platforms for infrastructure and services is a defining trend in modern software development.<sup>9</sup> Demonstrating experience with at least one major cloud provider, such as Amazon Web Services (AWS), Google Cloud Platform (GCP), or Microsoft Azure, and their associated services (e.g., compute, storage, databases, serverless functions) is a substantial advantage for candidates targeting FAANG companies. These organizations operate extensively in the cloud, leveraging its scalability and flexibility. Projects that showcase the ability to build and deploy applications using cloud-native principles and services highlight a developer's readiness for their operational environment.

### **DevOps Practices (CI/CD)**

Efficient software development and deployment in large organizations necessitate the adoption of DevOps practices, particularly continuous integration and continuous delivery (CI/CD) pipelines.<sup>3</sup> Experience in setting up and utilizing CI/CD pipelines with popular

tools like Jenkins, GitHub Actions, or similar platforms demonstrates an understanding of automation and efficient development workflows. FAANG companies highly value speed and efficiency in their development processes, and projects that incorporate well-implemented CI/CD pipelines showcase a developer's commitment to these practices.

### **Open Source Contributions**

Meaningful contributions to reputable open-source projects can significantly enhance a candidate's profile.<sup>56</sup> FAANG companies often contribute to and rely on open-source software, and active participation in these communities demonstrates a developer's technical skills, collaboration abilities, and genuine passion for technology. It provides tangible evidence of their capacity to work within established codebases, contribute effectively to team efforts, and engage with the broader development community.

## **3. Project Idea 1: Building a Scalable E-commerce Microservices Platform**

Developing a comprehensive e-commerce platform using a microservices architecture provides an excellent opportunity to showcase a wide range of skills highly relevant to FAANG companies. This project allows for the demonstration of expertise in designing, building, and deploying a complex distributed system.

### **Detailed Outline of Core Services**

- **Product Catalog Service:** This service is responsible for managing all product-related information, including the product's name, description, price, images, categories, and current inventory levels. It should expose API endpoints for creating new products, retrieving product details (individually or by category), updating existing product information, and deleting products. The underlying data can be stored in a structured database like PostgreSQL or MySQL <sup>5</sup>, with the potential to use a NoSQL database like MongoDB for more flexible product attributes if needed. Implementing this service demonstrates proficiency in database interaction, API design principles (potentially using RESTful APIs with JSON payloads), and effective data modeling. Managing product information is a fundamental aspect of any e-commerce platform, and showcasing the ability to handle this efficiently is crucial.
- **Order Management Service:** This service handles the entire lifecycle of an order, from placement to delivery and returns. Its functionalities include processing new order placements, tracking the status of each order, managing shipment details, and handling product returns. The service should provide API endpoints for creating new orders, updating the status of an order (e.g., processing, shipped, delivered), and retrieving detailed information about a specific order. A relational database like PostgreSQL or MySQL <sup>5</sup> would be suitable for storing order data. Building this service demonstrates an understanding of transaction management in a distributed environment and the ability to manage stateful operations across multiple services. Order processing is central to the business logic of an e-commerce platform.
- **User Authentication and Authorization Service:** Securing the platform is paramount, and this service manages user registration, login processes, and user roles. It should provide API endpoints for user registration, user login (generating authentication tokens), and role-based access control to protect other services and functionalities. Implementing security using Spring Security with JWT (JSON Web Tokens) is a strong choice.

Integration with an identity provider like Okta [9](#) can further enhance the project by demonstrating experience with modern authentication and authorization mechanisms. This service showcases an understanding of security best practices and identity management in a distributed system.[5](#) Protecting user data and ensuring secure access are essential considerations for any online platform.

- **Shopping Cart Service:** This service is responsible for managing individual user shopping carts. Functionalities include allowing users to add items to their cart, remove items, and calculate the total cost of the items in the cart. For optimal performance, especially with frequent updates, an in-memory data store like Redis [73](#) is a suitable choice for storing shopping cart data. This demonstrates the ability to use caching mechanisms and efficient data structures to optimize performance.[6](#) Providing a smooth and responsive shopping experience is crucial for user satisfaction.
- **Payment Processing Service (Simplified Mock):** While a full integration with a real payment gateway might be complex, a simplified mock implementation can still demonstrate understanding of the payment processing flow. This service would handle the initiation of payments and the management of payment status. API endpoints could include processing payment requests and updating payment status. Even a dummy implementation or integration with a sandbox environment of a payment gateway can showcase understanding of integration with external services and the handling of financial transactions, albeit in a simplified manner.[33](#) Processing payments is a critical component of any e-commerce platform.

### Technology Stack Suggestions

- **Backend:** Utilize Java 11 or a later version along with the Spring Boot 3.x framework.[3](#) Spring Boot's rapid development capabilities and extensive ecosystem make it well-suited for building microservices.
- **Containerization:** Employ Docker for containerizing each microservice [all microservices snippets]. This ensures portability and consistent deployment across different environments.
- **Orchestration:** For local development and testing, Kubernetes using minikube or Docker Compose [3](#) can be used. For demonstrating deployment to a more realistic environment, consider using a cloud-based Kubernetes service (like EKS, GKE, or AKS).
- **Database:** Choose either PostgreSQL [5](#) or MySQL [4](#) as the primary database for services like Product Catalog and Order Management.
- **API Gateway:** Implement an API Gateway using Spring Cloud Gateway.[1](#) This acts as a single entry point for all client requests, handling routing, security, and potentially rate limiting.
- **Service Discovery:** Use Netflix Eureka [9](#) for service discovery, allowing services to dynamically locate and communicate with each other.
- **Inter-service Communication:** Implement inter-service communication using either REST over HTTP for synchronous communication or an asynchronous messaging system like RabbitMQ [1](#) or Kafka [22](#) for more loosely coupled interactions. The choice depends on the specific requirements of each interaction.

The selection of these technologies reflects familiarity with industry-standard tools and frameworks commonly used in building scalable microservices architectures. Utilizing a modern and relevant technology stack demonstrates up-to-date knowledge and adherence to best practices.

## Key Implementation Steps

The implementation should involve a detailed breakdown of each service, including defining clear API contracts using OpenAPI/Swagger for documentation and contract-first development. Focus on creating well-defined data models for each entity (e.g., Product, Order, User, Cart Item) and implementing the core business logic within each service. Emphasize modularity in the codebase to ensure each service can be developed, tested, and deployed independently. Comprehensive unit and integration tests are crucial to ensure the reliability of each service.

## Dockerization Strategy

For each microservice, create a Dockerfile that specifies the base image (e.g., an official OpenJDK image **32**), copies the application code, installs any necessary dependencies, and defines the command to start the service. Optimize Docker image sizes by using multi-stage builds **35** to separate the build environment from the runtime environment, resulting in smaller and more efficient images. Define clear container startup commands and expose the necessary ports for each service.

## Considerations for Scalability and Resilience

To demonstrate an understanding of scalability, implement horizontal scaling capabilities for each service using Kubernetes or Docker Swarm. This involves configuring the orchestration platform to run multiple instances of each service and distribute traffic among them. Implement load balancing at the API Gateway level to evenly distribute incoming requests across the available service instances. Utilize service discovery to allow services to dynamically find and communicate with each other, especially as the number of instances scales up or down. For resilience, consider implementing circuit breaker patterns using a library like Resilience4j (as mentioned in **12**) to prevent cascading failures between services and improve the overall fault tolerance of the platform.

## 4. Project Idea 2: Developing a Distributed Task Queue with Java and Docker

Building a distributed task queue using Java and Docker is another excellent project idea that showcases understanding of distributed systems and concurrency. This type of system is fundamental to many large-scale applications for handling background processing and asynchronous tasks.

### Outline of the Architecture

The architecture of a distributed task queue typically involves three main components:

- **Producers:** These are applications or services that create tasks and enqueue them into the task queue.
- **Broker:** This is the central component that holds the tasks and ensures reliable delivery to consumers. Examples include Redis, RabbitMQ, and Kafka.
- **Consumers:** These are applications or services that dequeue tasks from the broker and process them.

### Technology Choices



- **Backend:** Use Java with either Project Loom [74](#) to explore lightweight concurrency with virtual threads, or utilize traditional Java concurrency mechanisms like threads and executors. Project Loom offers a modern approach to handling concurrency that can be highly efficient for I/O-bound tasks.
  - **Containerization:** Employ Docker to containerize the producer, broker (if self-hosted), and consumer components.
  - **Broker:** Choose one of the following popular message brokers:
  - **Redis:** A versatile in-memory data store that can be used as a message broker with libraries like Jedis or Lettuce.[73](#) Redis is known for its speed and simplicity.
  - **RabbitMQ:** A robust and widely used message broker that supports various messaging patterns.[1](#) It offers features like message durability and complex routing.
  - **Kafka:** A distributed streaming platform suitable for high-throughput message processing.[22](#) Kafka is designed for building real-time data pipelines and streaming applications.
- Offering these choices allows flexibility for the user to tailor the project to their specific interests and showcase different messaging paradigms and their associated trade-offs.

### Detailed Implementation Plan

1. **Defining the Task Structure:** Create a Java class that represents a task. This class should include relevant attributes depending on the type of work the task will perform. For example, if the task involves processing an image, the task class might include the image file path or URL and desired processing parameters.
2. **Implementing Producer Logic:** Develop the producer component, which will be responsible for creating task instances and enqueueing them to the chosen message broker. Use the appropriate client library for the selected broker (e.g., Jedis or Lettuce for Redis, the RabbitMQ Java client, or the Kafka client).
3. **Developing the Broker Component:**
  - **Redis:** If using Redis, implement the queue management logic using Redis commands like LPUSH (to add to the queue) and BRPOP (to retrieve from the queue with blocking).
  - **RabbitMQ/Kafka:** If using RabbitMQ or Kafka, configure the necessary queues or topics and exchanges (in RabbitMQ) to route messages appropriately.
4. **Implementing Consumer Logic:** Develop the consumer component(s) that will dequeue tasks from the broker and process them. Implement concurrent processing to handle multiple tasks simultaneously. Consider using Java's `ExecutorService` or, if using Project Loom, simply create multiple virtual threads. Ensure that the consumer logic handles potential exceptions gracefully and implements idempotency if task processing should be safe to retry.
5. **Adding Advanced Features:** Implement mechanisms for task acknowledgment (confirming successful processing), retries with exponential backoff for transient failures, and dead-letter queues to handle tasks that fail repeatedly. These features demonstrate a deeper understanding of building reliable distributed systems.

### Docker Setup

Create Dockerfiles for each component of the task queue (producer, consumer, and the broker if you choose to self-host it within Docker). Then, create a Docker Compose file to

define and run all the components locally. This file should specify the dependencies between the components and configure network settings as needed.

### Focus on Concepts

This project should emphasize key distributed systems concepts such as message durability (if using RabbitMQ or Kafka, configure persistence to ensure messages are not lost in case of broker failures), fault tolerance (design the consumers to handle broker outages and their own failures gracefully), handling backpressure (implement mechanisms to prevent producers from overwhelming consumers, perhaps by limiting the queue size or using rate limiting), and the ability to scale the number of consumers based on the workload. These considerations are highly relevant to the challenges faced in FAANG environments.

## 5. Project Idea 3: Creating a Simple Cloud-Native Application on a Public Cloud Platform (e.g., AWS, GCP, Azure)

Developing a simple yet functional application that leverages the services of a major public cloud platform is an excellent way to demonstrate cloud-native development skills, which are highly valued by FAANG companies. Choosing one of the major providers—AWS, GCP, or Azure—and utilizing their platform-specific services will showcase practical experience in a cloud environment.

### Outline of a Practical Application

Here are a few options for practical cloud-native applications:

- **Option 1 (AWS): A RESTful API for image resizing.** Users can upload images to an Amazon S3 bucket. This upload event can trigger an AWS Lambda function, written in Java, to automatically resize the image. The resized image can then be stored back in another S3 bucket. The API for uploading images and retrieving resized images can be exposed through Amazon API Gateway. This option demonstrates the use of serverless computing (Lambda), object storage (S3), and API management (API Gateway).
- **Option 2 (GCP): A simple note-taking web application.** The backend of this application, built using Spring Boot and containerized with Docker, can be deployed to Google Cloud Run, a fully managed serverless container execution service. Notes can be stored in Google Cloud Firestore, a NoSQL document database. This showcases containerized application deployment on a serverless platform (Cloud Run) and the use of a cloud-based NoSQL database (Firestore).
- **Option 3 (Azure): A data transformation API.** Users can send data (e.g., in JSON format) to an Azure Function, written in Java. This function can perform a specific data transformation (e.g., converting units, filtering data) and store the transformed data in Azure Blob Storage (for unstructured data) or Azure Cosmos DB (for structured NoSQL data). This option highlights the use of serverless functions (Azure Functions) and Azure's storage and database services.

Choosing one of these options provides a concrete scenario for developing a cloud-native application.

### Leveraging Cloud-Specific Services

The implementation should deeply leverage the specific services offered by the chosen cloud platform. This includes detailed explanations of how to set up and configure these services. For example, in the AWS image resizing application, the report should detail how to create S3 buckets, configure Lambda function triggers, write the Java code for the Lambda function using the AWS SDK for Java, and set up the API Gateway with appropriate endpoints and integrations. Similarly, for the GCP option, it should cover how to containerize the Spring Boot application with Docker, deploy it to Cloud Run, and interact with Cloud Firestore using the appropriate GCP client libraries. For the Azure option, it should describe how to create an Azure Function, write the Java code for data transformation, and use the Azure SDK for Java to interact with Blob Storage or Cosmos DB.

### **Detailed Steps for Building the Java Backend**

Regardless of the chosen cloud platform, the Java backend component will likely involve using a framework like Spring Boot to create the application logic. This includes defining API endpoints (if applicable), implementing the core functionality (image resizing, note storage, data transformation), and handling data persistence using the cloud-specific database or storage services. The implementation should follow best practices for Java development, including proper error handling, logging, and test coverage.

### **Dockerizing the Application for Cloud Deployment**

If the chosen application involves containerization (as in the GCP option, or if you choose to containerize the other options as well), create an optimized Docker image for the cloud environment. Consider using a lightweight base image, such as an official OpenJDK slim image from Docker Hub **32**, to minimize the image size. Ensure that the Dockerfile includes all necessary dependencies and configurations for the application to run correctly in the cloud environment.

### **Instructions on Deploying and Managing the Application**

Provide a step-by-step guide on how to deploy the application to the chosen cloud platform. This should include using the platform's command-line interface (e.g., AWS CLI, gcloud CLI, Azure CLI) or the web console. Detail how to set up necessary environment variables, configure monitoring using the cloud provider's monitoring services (e.g., AWS CloudWatch, Google Cloud Monitoring, Azure Monitor), and perform basic scaling of the application if the cloud service supports it (e.g., configuring auto-scaling for Lambda or Cloud Run).

## **6. Showcasing Your Projects on GitHub**

Once a project is implemented, effectively showcasing it on GitHub is crucial for making a positive impression on recruiters and hiring managers. Several best practices should be followed:

- **Repository Organization and Code Quality:** Organize the project repository logically with clear directory structures. Adhere to consistent Java coding standards, potentially using a tool like Checkstyle **39** to ensure code style consistency. Write clean, well-commented code that is easy to understand and maintain. Include build scripts (e.g., Maven or Gradle) to allow others to easily build and run the project.



- **Compelling README Files:** The README file is the first point of contact for anyone visiting the repository. It should clearly articulate the motivation behind the project, the goals it aims to achieve, and the key features and functionalities it offers. Detail the technology stack used, providing a concise overview of each technology and its role in the project. Include comprehensive step-by-step instructions on how to set up the development environment and run the application. If the project involves an API, provide clear API documentation, ideally using OpenAPI/Swagger specifications. A well-written README tells the story of the project and makes it easy for others to understand and appreciate the work done.
- **Utilizing GitHub Features:** Leverage GitHub features like issues, pull requests, and project boards to showcase the development workflow. Even if working on the project individually, using issues to track tasks and pull requests to organize changes can demonstrate familiarity with standard development practices. If the project evolves over time, maintaining a clear history of commits and releases further enhances its value as a portfolio piece. Consider adding a license to the repository to specify how others can use the code.

## 7. Contributing to Open Source Projects

Contributing to open-source projects is a valuable way to demonstrate technical skills, collaboration abilities, and a genuine passion for technology, all of which are highly regarded by FAANG companies.<sup>61</sup> Meaningful contributions to well-known projects can significantly strengthen a FAANG application.

### Strategies for Finding Relevant Projects

Identify open-source projects that align with your interests and skills. Consider projects related to Java, Docker, microservices, distributed systems, or cloud-native technologies. Table 2 provides a list of relevant open-source projects identified from the research material, including Apache Airavata <sup>24</sup>, Elasticsearch <sup>56</sup>, Jetty <sup>75</sup>, RxJava <sup>75</sup>, Apache Cassandra <sup>22</sup>, Apache HBase <sup>22</sup>, Apache Zookeeper <sup>22</sup>, Atomix <sup>22</sup>, Trino <sup>56</sup>, OpenMetadata <sup>56</sup>, QuestDB <sup>56</sup>, Strongbox <sup>56</sup>, conductor-oss/conductor <sup>39</sup>, and docker-java/docker-java.<sup>40</sup> Explore platforms like GitHub, GitLab, and Bitbucket to discover projects in these domains.

### Guidance on Making Meaningful Contributions

Start by looking for issues labeled "good first issue" <sup>56</sup> as these are typically designed for newcomers to the project. Meaningful contributions can include fixing bugs, implementing small enhancements, improving documentation, or writing tests. Focus on areas where your skills and interests align with the project's needs.

### Best Practices for Submitting Pull Requests

Before submitting a pull request, ensure that your code adheres to the project's coding standards and that all tests pass. Write clear and concise commit messages that explain the changes made. In your pull request description, clearly describe the problem you are addressing and the solution you have implemented. Be respectful and constructive in your interactions with the project maintainers and be prepared to address any feedback they provide.

### Emphasizing the Benefits for FAANG Applications

Highlight any open-source contributions on your resume and GitHub profile. In interviews, be prepared to discuss your contributions, the challenges you faced, and what you learned from the experience. Contributing to open source demonstrates practical coding skills, the ability to work collaboratively within a team, familiarity with development workflows (like Git and pull requests), and a genuine passion for technology, all of which are highly valued by FAANG recruiters.

## 8. Conclusion

Building a strong GitHub portfolio with well-implemented projects is a highly effective strategy for mid-senior Java/Docker developers looking to break into FAANG companies. The project ideas outlined in this report—building a scalable e-commerce microservices platform, developing a distributed task queue, or creating a cloud-native application on a public cloud—offer significant opportunities to showcase the key skills and knowledge that FAANG recruiters seek. When selecting a project, carefully consider your interests and the amount of time you can dedicate to it. Remember that thorough implementation, clear and comprehensive documentation on GitHub, and active engagement with the development community, including potential contributions to relevant open-source projects, will significantly enhance your profile and increase your chances of success in the competitive FAANG hiring landscape.

**Table 1: Comparison of Project Ideas**

| Feature              | Scalable E-commerce Platform  | Distributed Task Queue                                       | Cloud-Native Application   |
|----------------------|---|--|--|
| Core Focus           | Microservices, Scalability, API Design, Data Management   | Distributed Systems, Concurrency, Message Queuing            | Cloud Platforms, Serverless, Cloud Services Integration                |
| Technologies         | Java, Spring Boot, Docker, Kubernetes/Compose, Databases, API Gateway, Service Discovery, Messaging | Java, Project Loom/Concurrency, Docker, Redis/RabbitMQ/Kafka | Java, Spring Boot/Framework, Docker (Optional), AWS/GCP/Azure Services |
| Skills Demonstrated  | Full-stack (backend), System Design, Distributed Architecture, Security                             | Concurrency, Distributed Systems, Messaging, Fault Tolerance | Cloud Computing, Serverless, Platform-Specific Services, Deployment    |
| Complexity           | High  | Medium to High   | Medium   |
| Real-World Relevance | Very High (E-commerce is a common and complex domain)   | High (Task queues are fundamental to many applications)      | High (Cloud-native development is increasingly important)              |
| Showcase Potential   | Wide range of skills, complex architecture  | Deep understanding of distributed systems and concurrency    | Practical experience with a specific cloud platform                    |

**Table 2: Relevant Open Source Projects from Research**

| <b>Project Name</b>         | <b>Description</b>   | <b>Relevant Snippet(s)</b> |
|-----------------------------|--|----------------------------|
| Apache Airavata             | Software framework for executing and managing computational jobs on distributed computing resources.           | <b>24</b>                  |
| Elasticsearch               | Open Source, Distributed, RESTful Search Engine.   | <b>56</b>                  |
| jetty / jetty.project       | Eclipse Jetty® - Web Container & Clients - supports HTTP/2, HTTP/1.1, HTTP/1.0, websocket, servlets, and more. | <b>75</b>                  |
| ReactiveX / RxJava          | Reactive Extensions for the JVM – a library for composing asynchronous and event-based programs.               | <b>75</b>                  |
| Apache Cassandra            | A highly-scalable partitioned row store (NoSQL database).  | <b>22</b>                  |
| Apache HBase                | A Hadoop database, a distributed, scalable, big data store.  | <b>22</b>                  |
| Apache Zookeeper            | Highly reliable distributed coordination service.  | <b>22</b>                  |
| atomix                      | Fully featured framework for building fault-tolerant distributed systems.                                      | <b>22</b>                  |
| Trino (formerly Presto SQL) | A distributed SQL query engine for big data.   | <b>56</b>                  |
| OpenMetadata                | An all-in-one platform for data discovery, data quality, observability, governance, data lineage.              | <b>56</b>                  |
| QuestDB                     | A fast open source SQL time series database.   | <b>56</b>                  |
| Strongbox                   | An artifact repository manager written in Java.  | <b>56</b>                  |
| conductor-oss/ conductor    | An event driven orchestration platform providing durable and highly resilient execution engine.                | <b>39</b>                  |
| docker-java/ docker-java    | Java client library for the Docker Remote API.   | <b>40</b>                  |